SAYAN GHOSH

Washington State University ◇ Pullman, WA 99163
sayan.ghosh@wsu.edu ◇ https://sg0.github.io

EDUCATION

- **January 2015 - Present**
  PhD, Computer Science,
  Washington State University, School of EECS, Pullman, WA.
  Adviser: Dr. Assefaw Gebremedhin
  Expected graduation: Spring, 2019
  Thesis title: *Supporting Asynchronous Parallel Programming Models for Scientific and Irregular Applications*
  PhD program commitee members: Drs. Assefaw Gebremedhin (WSU), Carl Hauser (WSU), Ananth Kalyanaraman (WSU), Pavan Balaji (ANL) and Mahantesh Halappanavar (PNNL).
- **2012 - 2014**
  PhD studies, Computer Science,
  University of Houston
  Adviser: Dr. Barbara Chapman
- **2010 - 2012**
  Master of Science (Thesis), Computer Science,
  University of Houston, Houston, TX.
  Graduated: August 2012
  Thesis title: *Energy Efficiency of Parallel Scientific Kernels*
  Adviser: Dr. Barbara Chapman
- **2002 - 2006**
  Bachelor of Technology, Information Technology
  Asansol Engineering College, Asansol, India.
  Graduated: July 2006

EXPERIENCES

Washington State University, Pullman, WA Jan 2015-Present
Graduate Research Assistant *Advisor: Dr. Assefaw Gebremedhin*
· *Focus*: Graph analytics, Combinatorial algorithms, One-sided programming models.

University of Houston, Houston, TX Jan 2011-Dec 2014
Graduate Research Assistant *Advisor: Dr. Barbara Chapman*
· *Focus*: Power/energy analysis and modeling of scientific kernels, Application parallelization using compiler directives, One-sided programming models.

University of Texas Health Science Center, Houston, TX Jan-Dec 2010
Graduate Research Assistant *Advisor: Dr. Stefan Birmanns*
· *Focus*: Application parallelization using compiler directives.

Thomson Reuters, Bangalore, India Jul 2008-Dec 2009
Software Engineer
· *Focus*: Database design and development.

NTT Innovation Institute, Inc., Bangalore, India Jul 2006-Jul 2008
Software Engineer
· *Focus*: Database design and development.

PUBLICATIONS /PRESENTATIONS

- [**Journal**] Sayan Ghosh, Terrence Liao, Henri Calandra and Barbara Chapman. *Performance of CPU/GPU compiler directives on ISO/TTI kernels.* Computing Journal, Springer Vienna (2013).
- **Conferences**
  - Sayan Ghosh, Mahantesh Halappanavar, Antonino Tumeo, Ananth Kalyanaraman, Assefaw Gebremedhin. *Scalable Distributed Memory Community Detection Using Vite.* 22nd IEEE High Performance Extreme Computing Conference (HPEC 2018). (*Student Innovation Award*)
  - Sayan Ghosh, Mahantesh Halappanavar, Antonino Tumeo, Ananth Kalyanaraman, Hao Lu, Daniel Chavarrià-Miranda, Arif Khan, Assefaw Gebremedhin. *Distributed Louvain Algorithm for Graph Community Detection.* 32nd IEEE International Parallel and Distributed Processing Symposium (IPDPS 2018).
  - Sayan Ghosh, Assefaw Gebremedhin. *Parallelization of Bin Packing on Multicore Systems.* 23rd International Conference on High Performance Computing, Data, and Analytics (HiPC 2016).
  - Sayan Ghosh, Jeff Hammond, Antonio J. Peña, Pavan Balaji, Assefaw Gebremedhin, Barbara Chapman. *One-Sided Interface for Matrix Operations using MPI-3 RMA: A Case Study with Elemental.* 45th International Conference on Parallel Processing (ICPP 2016).
  - Naveen Namashivayam, Sayan Ghosh, Dounia Khaldi, Deepak Eachempati, Barbara Chapman. *Native Mode-Based Optimizations of Remote Memory Accesses in OpenSHMEM for Intel Xeon Phi.* 8th International Conference on Partitioned Global Address Space Programming Models (PGAS 2014). (*Best Paper*)
- **Workshops**
  - Sayan Ghosh, Mahantesh Halappanavar, Antonino Tumeo, Ananth Kalyanaraman, Assefaw Gebremedhin. *miniVite: A Graph Analytics Benchmarking Tool for Massively Parallel Systems.* Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS 2019).
  - Priyanka Ghosh, Jeff Hammond, Sayan Ghosh, Barbara Chapman. *Performance Analysis of the NWChem TCE for Different Communication Patterns.* Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS 2013).
  - Jeff Hammond, Sayan Ghosh, Barbara Chapman. *Implementing OpenSHMEM using MPI-3 one-sided communication.* 1st OpenSHMEM Workshop: Experiences, Implementations and Tools (2013).
  - Sayan Ghosh, Sunita Chandrasekaran, Barbara Chapman. *Statistical modeling of power/energy of scientific kernels on a multi-GPU system.* Power Measurement and Profiling Workshop (PMP), in conjunction with International Green Computing Conference (IGCC 2013).
  - Sayan Ghosh, Terrence Liao, Henri Calandra, Barbara Chapman. *Experiences with OpenMP, PGI, HMPP and OpenACC directives on ISO/TTI kernels.* 5th International Workshop on Multi/Manycore Computing Systems (MuCoCoS 2012).
  - Sayan Ghosh, Sunita Chandrasekaran, Barbara Chapman. *Energy Analysis of Parallel Scientific Kernels on Multiple GPUs.* Symposium of Application Accelerators in High Performance Computing (SAAHPC 2012).
- **Posters**
  - Sayan Ghosh, Assefaw Gebremedhin. *Towards a More Asynchronous GraphBLAS.* SIAM workshop on Combinatorial Scientific Computing (CSC 2016).
  - Sayan Ghosh, Sunita Chandrasekaran, Barbara Chapman. *Power and Energy Prediction of Multi-GPU kernels Using Non-linear Regression.* Nvidia GPU Technology Conference (GTC 2013).
  - Sayan Ghosh, Terrence Liao, Henri Calandra, Barbara Chapman. *Performance of ISO/TTI kernels on CPU/GPU using OpenMP, PGI, HMPP and OpenACC directives.* Rice Oil and Gas HPC Workshop (OGHPC 2013).
  - Sayan Ghosh, Sunita Chandrasekaran, Barbara Chapman. *Statistical Power and Energy Modeling of multi-GPU kernels.* General poster, International Conference for High Performance Computing, Networking, Storage, and Analysis (SC 2012).
  - Sayan Ghosh, Barbara Chapman. *Programming Strategies for GPUs and their Power Consumption.* General poster, International Conference on Parallel Architectures and Compilation Techniques (PACT 2012).

INTERNSHIPS

Pacific Northwest National Laboratory, Richland, WA May-Aug 2018
Supervisor: Drs. Mahantesh Halappanavar and Arif Khan

Focus: Distributed-memory graph analytic algorithms, such as community detection and maximal weight matching.

Pacific Northwest National Laboratory, Richland, WA May-Aug 2017
Supervisor: Dr. Mahantesh Halappanavar

Focus: Distributed-memory network community detection.

Argonne National Laboratory, Chicago, IL May-Aug 2016
Supervisors: Drs. Pavan Balaji and Yanfei Guo

Focus: C++ bindings to MPI-3 RMA.

Argonne National Laboratory, Chicago, IL May-Aug 2014
Supervisors: Drs. Pavan Balaji and Antonio J. Peña

Focus: Asynchronous interface for updating distributed matrices in Elemental, a distributed-memory dense linear algebra library.

Argonne National Laboratory, Chicago, IL May-Aug 2013
Supervisor: Dr. Jeff Hammond

Focus: Design and prototype of a one-sided communication runtime on top of MPI-3, that led to development of an OpenSHMEM implementation over MPI-3 RMA.

Total R&T, Houston, TX May-Aug 2012
Supervisors: Drs. Terrence Liao and Henri Calandra

Focus: Evaluation of directive based programming models like OpenMP, PGI, HMPP and OpenACC on Finite Difference kernels used in Oil and Gas exploration, on GPU and multicore CPUs.

Pacific Northwest National Laboratory, Richland, WA Jun-Sept 2011
Supervisors: Drs. Darren Kerbyson, Kevin Barker and Abhinav Vishnu

Focus: Power/energy profiling of scientific kernels on a multi-GPU platform.

ACTIVITIES/TRAVEL GRANTS

- Participant, 2018 Argonne Training Program on Extreme-Scale Computing (ATPESC), July 29-August 10, St. Charles, IL
- NSF/IEEE TCPP Travel grant, 32rd International Parallel and Distributed Processing Symposium (IPDPS 2018), Vancouver, BC, Canada
- NSF/IEEE TCPP Travel grant, 23rd International Conference on High Performance Computing, Data, and Analytics (HiPC 2016), Hyderabad, India
- Student Volunteer, Supercomputing 2016, Salt Lake City, Utah
- Booth setup personnel, Gulf Coast Advanced Supercomputing (GCAS) booth, Supercomputing 2014, New Orleans, Louisiana
- Booth duty at Gulf Coast Advanced Supercomputing (GCAS) booth, Supercomputing 2013, Denver, Colorado
- Student Volunteer at Architectural Support for Programming Languages and Operating Systems (ASPLOS 2013) conference, Rice University, Houston, TX
- Co-taught a classroom session on OpenACC at Nvidia Global Technology Conference (GTC), San Jose, CA, 2013
- Booth duty at OpenMP booth and Gulf Coast Advanced Supercomputing (GCAS) booth, Supercomputing 2012, Salt Lake City, UT
- Represented University of Houston in OpenMP booth at Multicore Developers Conference, San Jose, CA (2011 and 2012)

TEACHING ASSISTANCESHIPS

- Spring 2016, Washington State University, EECS, Distributed Computing, CPTS 464/564 (*Course Instructor:* Dr. Dave Bakken)
- Fall 2015, Washington State University, EECS, Computer Communication Networks, CPTS 455 (*Course Instructor:* Dr. Carl Hauser)
- Spring 2015, Washington State University, EECS, Distributed Computing, CPTS 464/564 (*Course Instructor:* Dr. Dave Bakken)
- Fall 2010, University of Texas Health Science Center, Introductory Course on Data Structures (*Course Instructor:* Dr. Stefan Birmanns). This was an unofficial appointment, just assisted my advisor in taking the course and prepared course materials.

SOFTWARE USAGE

- *Programming Languages:* C, C++, Python, CUDA, FORTRAN, R
- *Standards/Libraries:* MPI, OpenSHMEM, Global Arrays Toolkit
- *Directive Based API:* OpenMP, OpenACC
- *Profilers/Debuggers:* GNU GDB, HPCToolkit, Valgrind TAU, Intel VTune, CUDA Profiler
- *Version Control Systems:* SVN, Git, Mercurial
- *Tools/Packages:* Gnuplot, LaTeX, CMake, GNU Autoconf

MEMBERSHIPS

- ACM Special Interest Group in High Performance Computing (SIGHPC)
- Institute of Electrical and Electronics Engineers (IEEE), IEEE Computer Society
- Society for Industrial and Applied Mathematics (SIAM)